

# Facial Expression Classification Using PyTorch

Farid Vakili

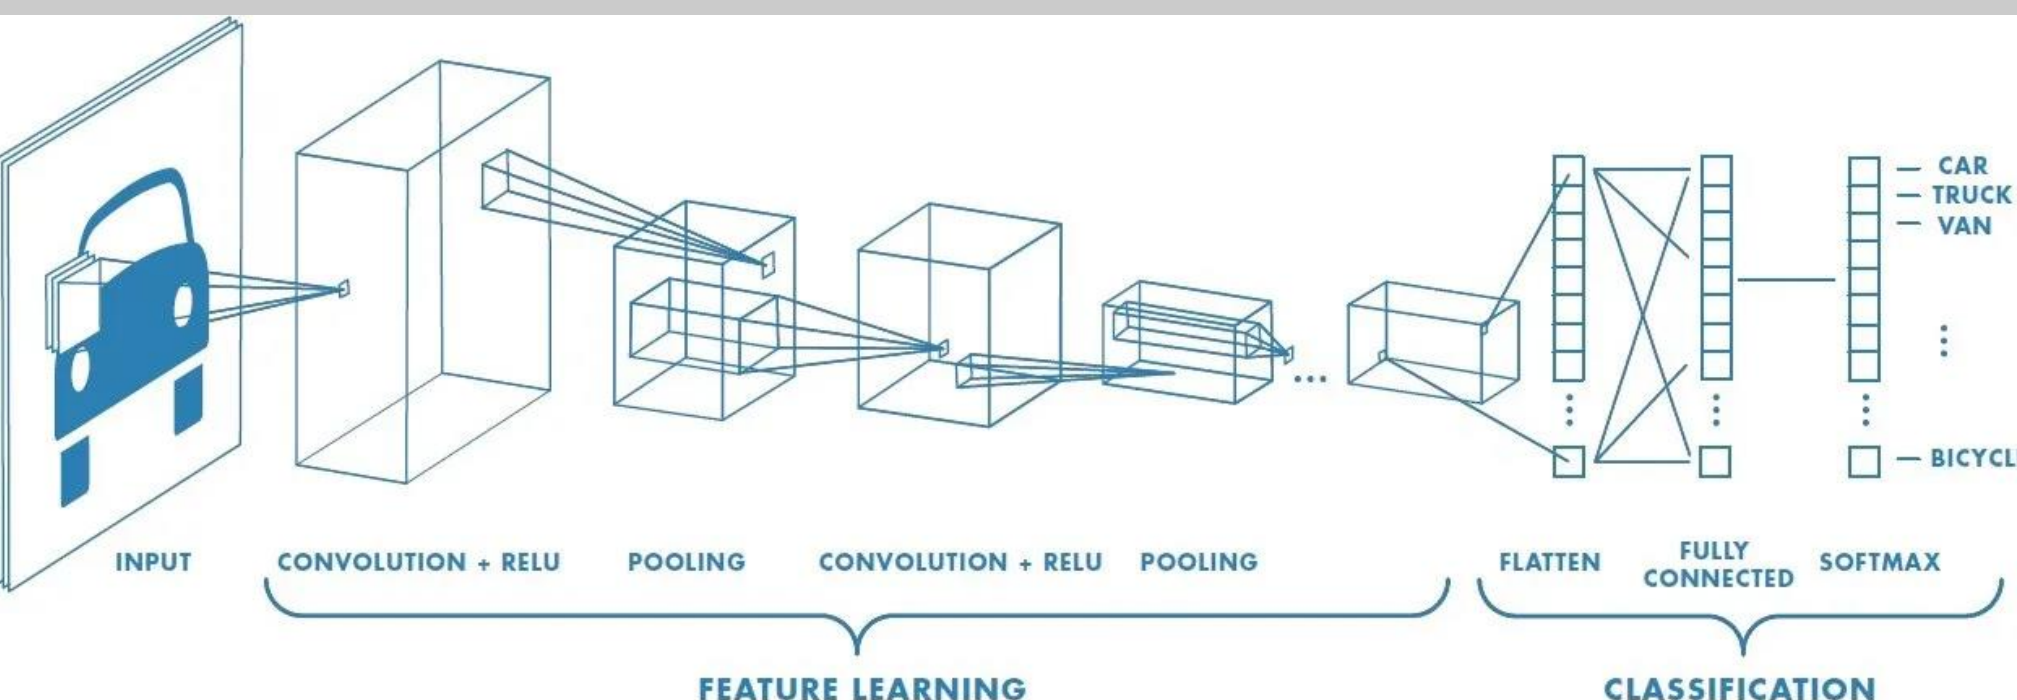
California State Polytechnic University Pomona

## ABSTRACT


Facial expression recognition is a crucial task in the field of computer vision with applications ranging from human-computer interaction to psychological studies. This project utilizes convolution neural network (CNN) to classify facial expressions using PyTorch (an open-source machine learning library)

## OBJECTIVES

Our objective was to implement a CNN for classifying facial expressions into three categories: Angry, Happy, and Neutral (we were also given a base line accuracy score to beat). The model was trained and evaluated on a publicly available dataset from a previous Kaggle competition



<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>



<https://www.linkedin.com/pulse/activation-functions-101-sigmoid-tanh-relu-softmax-ben-hammouda>

## MATERIALS & METHODS

### Dataset

The dataset from the "Challenges in Representation Learning: Facial Expression Recognition Challenge" on Kaggle <https://www.kaggle.com/c/challenges-in-representation-learning-facial-expression-recognition-challenge/data>

- **Training Set:** 16,175 grayscale images of size 48x48 pixels
- **Test Set:** 3,965 images of the same dimensions
- **Classes:** Angry (0), Happy (1), Neutral (2)

### Data Preprocessing

- **Normalization:** Pixel values scaled to the range [0, 1]
- **Reshaping:** Images reshaped to (48, 48, 1) to represent single-channel grayscale images

### Data Augmentation

To enhance model generalization, data augmentation techniques are applied:

- **Random Horizontal Flip:** Simulates mirror images
- **Random Rotation:** Images are rotated within a range of  $\pm 10$  degrees

### Model Architecture

The CNN model:

- **Convolutional Layers:** Three blocks with increasing filter sizes (32, 64, 128) and kernel size 3x3
  - Each block includes a convolutional layer, ReLU activation, batch normalization, and max pooling
- **Fully Connected Layers:**
  - A dense layer with 256 neurons and ReLU activation
  - Dropout layer to prevent overfitting
  - Output layer with 3 neurons (each class)

### Training Settings

- **Loss Function:** Cross-Entropy Loss
- **Optimizer:** Adam optimizer (initial learning rate 0.001)
- **Batch Size:** 64
- **Number of Epochs:** 20
- **Device:** Training uses GPU (if available)

### Validation Strategy

- **Train-Validation Split:** Split the training data to monitor validation performance
- **Early Stopping:** The model saves weights when the validation accuracy improves

## RESULTS

### Training and Validation Performance

Over 20 epochs, the model showed steady improvement in both training and validation metrics

• **Final Training Loss:** Decreased significantly (indicating good learning capacity)

• **Validation Accuracy:** Achieved approximately 80% accuracy

Epoch	[1/20]	Train Loss:	1.0169	Val Loss:	0.8143	Val Acc:	0.6357
Model saved!							
Epoch	[2/20]	Train Loss:	0.8081	Val Loss:	0.6818	Val Acc:	0.7008
Model saved!							
Epoch	[3/20]	Train Loss:	0.7150	Val Loss:	0.6412	Val Acc:	0.7237
Model saved!							
Epoch	[4/20]	Train Loss:	0.6641	Val Loss:	0.6075	Val Acc:	0.7365
Model saved!							
Epoch	[5/20]	Train Loss:	0.6248	Val Loss:	0.5874	Val Acc:	0.7480
Model saved!							
Epoch	[6/20]	Train Loss:	0.5958	Val Loss:	0.5709	Val Acc:	0.7694
Model saved!							
Epoch	[7/20]	Train Loss:	0.5794	Val Loss:	0.5605	Val Acc:	0.7655
Epoch	[8/20]	Train Loss:	0.5588	Val Loss:	0.5688	Val Acc:	0.7715
Model saved!							
Epoch	[9/20]	Train Loss:	0.5375	Val Loss:	0.5531	Val Acc:	0.7678
Epoch	[10/20]	Train Loss:	0.5253	Val Loss:	0.5586	Val Acc:	0.7814
Model saved!							
Epoch	[11/20]	Train Loss:	0.5128	Val Loss:	0.5386	Val Acc:	0.7762
Epoch	[12/20]	Train Loss:	0.4946	Val Loss:	0.5352	Val Acc:	0.7826
Model saved!							
Epoch	[13/20]	Train Loss:	0.4794	Val Loss:	0.5429	Val Acc:	0.7797
Epoch	[14/20]	Train Loss:	0.4591	Val Loss:	0.5416	Val Acc:	0.7826
Epoch	[15/20]	Train Loss:	0.4530	Val Loss:	0.5495	Val Acc:	0.7822
Epoch	[16/20]	Train Loss:	0.4356	Val Loss:	0.5664	Val Acc:	0.7797
...							
Epoch	[18/20]	Train Loss:	0.4109	Val Loss:	0.5426	Val Acc:	0.7923
Model saved!							
Epoch	[19/20]	Train Loss:	0.4037	Val Loss:	0.5512	Val Acc:	0.7915
Epoch	[20/20]	Train Loss:	0.3924	Val Loss:	0.5979	Val Acc:	0.7902

### Techniques to Improve Performance

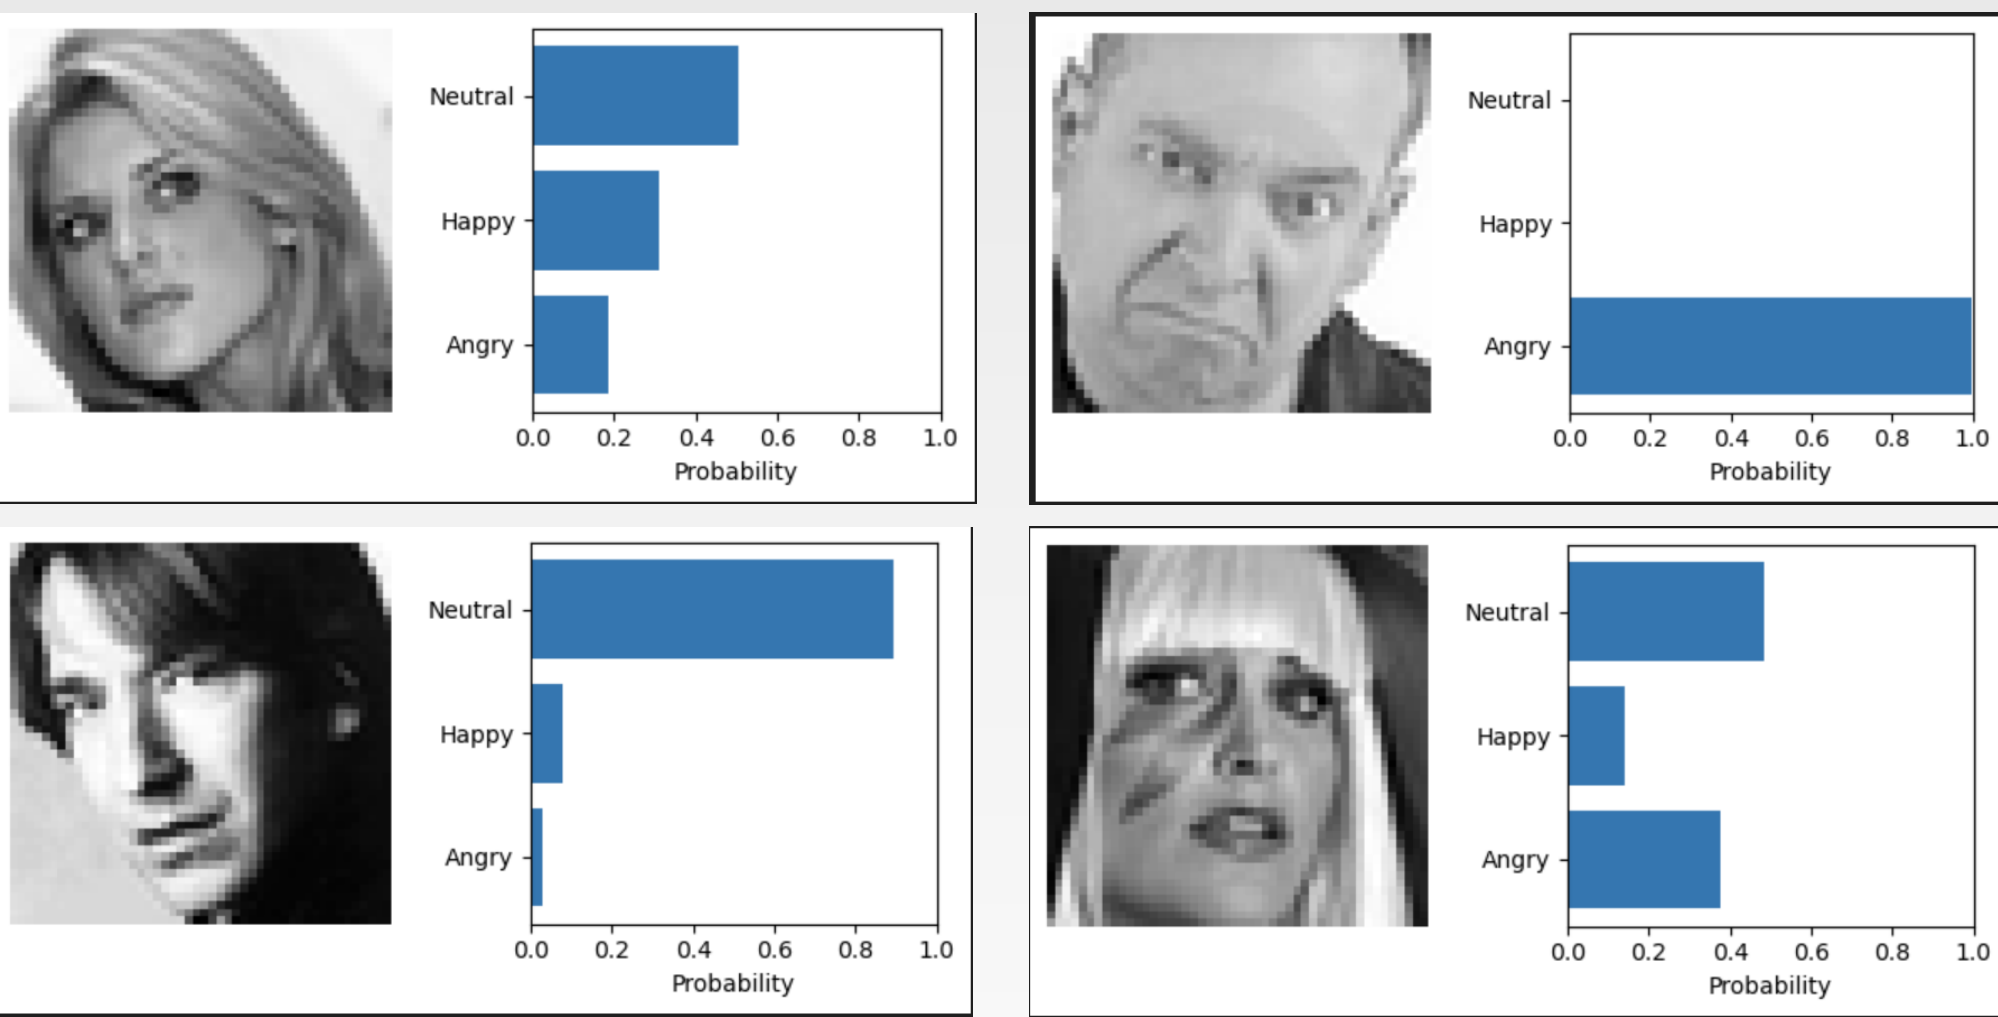
• **Data Augmentation:** Helped reduce overfitting by providing more varied training examples

• **Batch Normalization:** Improved training stability and convergence speed

• **Dropout Layer:** Prevented overfitting by randomly deactivating neurons during training

• **Hyperparameter Tuning:** Adjusted learning rates and batch sizes

### -Image Classification Probabilities-



## CONCLUSION

This assignment demonstrates the capability of CNNs in classifying facial expressions. Following the steps of the machine learning development process, our model generalizes pretty well to new data. Future research could explore deeper architectures, or incorporate attention mechanisms to further enhance performance

## REFERENCES & ACKNOWLEDGMENTS

1. Li, S., & Deng, W. (2020). "Deep Facial Expression Recognition: A Survey." *IEEE Transactions on Affective Computing*, 1-1
2. Alizadeh, Shima, and Azar Fazel. "Convolutional Neural Networks for Facial Expression Recognition" <https://Arxiv.Org/Abs/1704.06756>, 22 Apr. 2017, vision.stanford.edu/teaching/cs231n/reports/2016/pdfs/005\_Report.pdf

1. Andrew Ng – Stanford DeepLearning.AI <https://www.coursera.org/specializations/machine-learning-introduction>
2. Parth Dhameliya – Machine Learning Instructor Coursera <https://www.coursera.org/projects/facial-expression-recognition-with-pytorch> (demonstrated the function for image classification probability outputs)

